

COURSE 6

Relational Algebra

Relational Query Languages

- Query languages: Allow manipulation and **retrieval of data** from a database.
- Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- Query Languages **!=** programming languages!
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.

Formal Relational Query Languages

- Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:
 - Relational Algebra: More **operational**, very useful for representing execution plans.
 - Relational Calculus: Lets users describe what they want, rather than how to compute it. (**Non-operational**, declarative)

Preliminaries

- A query is applied to *relation instances*, and the result of a query is also a relation instance.
 - *Schemas of input* relations for a query are *fixed* (but query will run regardless of instance!)
 - The *schema for the result* of a given query is also *fixed*! Determined by definition of query language constructs.
- Positional vs. named-field notation:
 - Positional notation easier for formal definitions, named-field notation more readable.
 - Both used in SQL

Relational Algebra

■ Basic operations:

- Projection (π) Deletes unwanted columns from relation.
- Selection (σ) Selects a subset of rows from relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.

■ Additional operations:

- Intersection, join, division, renaming: Not essential, but (very!) useful.
- Since each operation returns a relation, **operations can be composed!** (Algebra is “closed”.)

Projection

- $L = (a_1, \dots, a_n)$ is a list of attributes (i.e. *a list of columns*) of the relation R
- Keeping vertical slices of a relation according to L

$$\pi_L(R) = \{ t \mid t_1 \in R \wedge \\ t.a_1 = t_1.a_1 \wedge \\ \dots \wedge \\ t.a_n = t_1.a_n \}$$

Projection (cont.)

$\pi_{cid, grade}(\text{Enrolled})$

$\pi_{cid, grade}(\$

<i>sid</i>	<i>cid</i>	<i>grade</i>
1234	Alg1	9
1235	Alg1	10
1234	DB1	10
1234	DB2	9
1236	DB1	7
1237	DB2	9
1237	DB1	5
1237	Alg1	10

) =

<i>cid</i>	<i>grade</i>
Alg1	9
Alg1	10
DB1	10
DB2	9
DB1	7
DB1	5

Projection (cont.)

Is $\pi_{\text{cid, grade}}(\text{Enrolled})$ equivalent to

SELECT cid, grade FROM Enrolled ?

NO! Relational algebra works with sets => no duplicates.

**SELECT DISTINCT cid, grade
FROM Enrolled**

Selection

- Selecting the t-uples of a relation **R** verifying a condition *c* (*selection predicate*).

$$\sigma_c(R) = \{ t \mid t \in R \wedge c \}$$

$$\sigma_{\text{grade} > 8}(\text{Enrolled}) = \{ t \mid t \in \text{Enrolled} \wedge \text{grade} > 8 \}$$

<i>sid</i>	<i>cid</i>	<i>grade</i>
1234	Alg1	9
1235	Alg1	10
1234	DB2	9
1236	DB1	7
1237	DB1	5
1237	Alg1	6

$$\sigma_{\text{grade} > 8} ($$

<i>sid</i>	<i>cid</i>	<i>grade</i>
1234	Alg1	9
1235	Alg1	10
1234	DB2	9

$$) =$$

Selection (cont.)

$\sigma_{\text{grade} > 8}(\text{Enrolled})$

```
SELECT DISTINCT *  
FROM Enrolled  
WHERE grade > 8
```

Confusing terms

Projection

$\pi_{\text{attr1}, \text{attr2}}(\text{Relation})$



SELECT DISTINCT attr₁, attr₂
FROM Relation



WHERE c

Selection $\sigma_c(\text{Relation})$

Selection Condition

- **Term Op Term** is a condition
 - where **Term** is an attribute name
 - or **Term** is a constant
 - **Op** is one of $<$, $>$, $=$, \neq etc.
- $(C1 \wedge C2)$, $(C1 \vee C2)$, $(\neg C1)$ are conditions where $C1$ and $C2$ are conditions

Composability

The result of an expression is a relation

$$\pi_{\text{cid, grade}}(\sigma_{\text{grade} > 8}(\text{Enrolled}))$$

$$\pi_{\text{cid, grade}}(\sigma_{\text{grade} > 8}(\begin{array}{|c|c|c|} \hline \textit{sid} & \textit{cid} & \textit{grade} \\ \hline 1234 & \text{Alg1} & 9 \\ 1235 & \text{Alg1} & 10 \\ 1234 & \text{DB1} & 10 \\ 1234 & \text{DB2} & 9 \\ 1236 & \text{DB1} & 7 \\ 1237 & \text{DB2} & 9 \\ 1237 & \text{DB1} & 5 \\ 1237 & \text{Alg1} & 10 \\ \hline \end{array})) =$$

<i>cid</i>	<i>grade</i>
Alg1	9
Alg1	10
DB1	10
DB2	9

Composability (cont.)

$$\pi_{\text{cid, grade}}(\sigma_{\text{grade} > 8}(\text{Enrolled}))$$

```
SELECT DISTINCT cid, grade  
FROM Enrolled  
WHERE grade > 8
```

$$\sigma_{\text{grade} > 8}(\pi_{\text{cid, grade}}(\text{Enrolled}))$$

What is the equivalent SQL query?

Can we always exchange the order of σ and π ?

Union, Intersection, Set-difference

- $R_1 \cup R_2 = \{ t \mid t \in R_1 \vee t \in R_2 \}$
- $R_1 \cap R_2 = \{ t \mid t \in R_1 \wedge t \in R_2 \}$
- $R_1 - R_2 = \{ t \mid t \in R_1 \wedge t \notin R_2 \}$

The relations R_1 and R_2 must be union compatible:

- same number of attributes (same *arity*)
- corresponding attributes have *compatible* domains and the *same name*

Set operations and equivalent SQL statements

$$R_1 \cup R_2$$

SELECT DISTINCT *
FROM R₁

UNION

SELECT DISTINCT *
FROM R₂

$$R_1 \cap R_2$$

SELECT DISTINCT *
FROM R₁

INTERSECT

SELECT DISTINCT *
FROM R₂

$$R_1 - R_2$$

SELECT DISTINCT *
FROM R₁

EXCEPT

SELECT DISTINCT *
FROM R₂

Are all operators essential?

$$R_1 \cap R_2 = ((R_1 \cup R_2) - (R_1 - R_2)) - (R_2 - R_1)$$

Compute all tuples
belonging to R_1 or R_2

Remove the ones
that
belong only to R_1

Remove the ones
that
belong only to R_2

Cartesian Product

- Combining two relations

$$R_1(a_1, \dots, a_n) \text{ and } R_2(b_1, \dots, b_m)$$

$$R_1 \times R_2 = \{ t \mid t_1 \in R_1 \wedge t_2 \in R_2 \\ \wedge t.a_1 = t_1.a_1 \dots \wedge t.a_n = t_1.a_n \\ \wedge t.b_1 = t_2.b_1 \dots \wedge t.b_m = t_2.b_m \}$$

SELECT DISTINCT *
FROM R₁, R₂

θ -Join

- Combining two relations R_1 and R_2 on a condition c

$$R_1 \otimes_c R_2 = \sigma_c (R_1 \times R_2)$$

Students $\otimes_{\text{Students.sid=Enrolled.sid}}$ Enrolled

```
SELECT DISTINCT *  
FROM Students, Enrolled  
WHERE Students.sid =  
Enrolled.sid
```

```
SELECT DISTINCT *  
FROM Students  
INNER JOIN Enrolled ON  
Students.sid=Enrolled.sid
```

The Equi-Join

- Combines two relations on a condition composed only of equalities of attributes of the first and second relation and projects only one of the redundant attributes (since they are equal)

$$R_1 \otimes_{E(c)} R_2$$

Courses

<i>cid</i>	<i>cname</i>
Alg1	Algorithms1
DB1	Databases1
DB2	Databases2

$$\otimes_{E(\text{Courses.cid} = \text{Enrolled.cid})}$$

Enrolled

<i>sid</i>	<i>cid</i>	<i>grade</i>
1234	Alg1	9
1235	Alg1	10
1234	DB1	10
1234	DB2	9
1236	DB1	7

=

<i>cname</i>	<i>sid</i>	<i>cid</i>	<i>grad</i>
Algorithms1	1234	Alg1	9
Algorithms1	1235	Alg1	10
Databases1	1234	DB1	10
Databases2	1234	DB2	9
Databases1	1236	DB1	7

The Natural Join

- Combines two relations on the equality of the attributes with the same names and projects only one of the redundant attributes

$$R_1 \otimes R_2$$

<i>Courses</i>			<i>Enrolled</i>			
<i>cid</i>	<i>cname</i>		<i>sid</i>	<i>cid</i>	<i>grade</i>	
Alg1	Algorithms1	⊗	1234	Alg1	9	
DB1	Databases1		1235	Alg1	10	
DB2	Databases2		1234	DB1	10	
			1234	DB2	9	
			1236	DB1	7	
		=				
			<i>cname</i>	<i>sid</i>	<i>cid</i>	<i>grad</i>
			Algorithms1	1234	Alg1	9
			Algorithms1	1235	Alg1	10
			Databases1	1234	DB1	10
			Databases2	1234	DB2	9
			Databases1	1236	DB1	7

Division

- Not supported as a primitive operator, but useful
- Let R_1 have 2 fields, x and y ; R_2 have only field y :

$$R_1 / R_2 = \{ \langle x \rangle \mid \exists \langle x, y \rangle \in R_1 \quad \forall \langle y \rangle \in R_2 \}$$

i.e., **R_1 / R_2 contains all x tuples such that for every y tuple in R_2 , there is an xy tuple in R_1 .**

Or: If the set of y values associated with an x value in R_1 contains all y values in R_2 , the x value is in R_1 / R_2 .

- In general, x and y can be any lists of fields; y is the list of fields in R_2 , and $x \cup y$ is the list of fields of R_1 .

Expressing R_1/R_2 Using Basic Operators

- Division is not essential op; just a useful shorthand.
 - (Also true of joins, but joins are so common that systems implement joins specially.)
- *Idea*: For R_1/R_2 , compute all x values that are not 'disqualified' by some y value in R_2 .
 - x value is *disqualified* if by attaching y value from R_2 , we obtain an xy tuple that is not in R_1 .

Disqualified x values: $\pi_x ((\pi_x(R_1) \times R_2) - R_1)$

$$R_1/R_2 = \pi_x(R_1) - \text{all disqualified values}$$

Renaming

- If attributes or relations have the same name (for instance when joining a relation with itself) it may be convenient to rename one

$$\rho(\mathbf{R}' (N_1 \rightarrow \mathbf{N}'_1, N_2 \rightarrow \mathbf{N}'_2), R)$$

alternative notation: $\rho_{\mathbf{R}' (N'_1, N'_2)}(R),$

- The new relation \mathbf{R}' has the same instance as R , its schema has attribute \mathbf{N}'_i instead of attribute N_i

Renaming (cont.)

$\rho(\text{Courses2 (cid} \rightarrow \text{code,}$
 $\text{cname} \rightarrow \text{description)},$
 $\text{Courses})$

Courses

<i>cid</i>	<i>cname</i>	<i>credits</i>
Alg1	Algorithms1	7
DB1	Databases1	6
DB2	Databases2	6



Courses2

<i>code</i>	<i>description</i>	<i>credits</i>
Alg1	Algorithms1	7
DB1	Databases1	6
DB2	Databases2	6

```
SELECT cid as code,  
       cname as description,  
       credits  
FROM Courses Courses2
```

Assignment Operation


- The assignment operation \leftarrow provides a convenient way to express complex queries.
 - Assignment must always be made to a temporary relation variable

$$\text{Temp} \leftarrow \pi_x(R_1 \times R_2)$$

- The result to the right of the \leftarrow is assigned to the relation variable on the left of the \leftarrow .
- May use variable in subsequent expressions.
 - $\text{result} \leftarrow \text{Temp} - R_3$

Complex Expression

$$R_1 \cup (R_2 \cap \pi_b (R_3 \times \rho(R_4(a \rightarrow b), R_5)))$$



Find names of students enrolled at 'BD1'

Solution 1: $\pi_{\text{name}} ((\sigma_{\text{cid}='BD1'}(\text{Enrolled})) \otimes \text{Students})$

Solution 2: $\rho (\text{Temp}_1, \sigma_{\text{cid}='BD1'}(\text{Enrolled}))$
 $\rho (\text{Temp}_2, \text{Temp}_1 \otimes \text{Students})$
 $\pi_{\text{name}} (\text{Temp}_2)$

Solution 3: $\pi_{\text{name}} (\sigma_{\text{cid}='BD1'}(\text{Enrolled} \otimes \text{Students}))$

Find names of students enrolled at a 5 credits course

- Information about course credits only available in Courses; so need an extra join:

$$\pi_{\text{name}} ((\sigma_{\text{credits}=5}(\text{Courses})) \otimes \text{Enrolled} \otimes \text{Students})$$

- A more efficient solution:

$$\pi_{\text{name}} (\pi_{\text{sid}} (\pi_{\text{cid}} (\sigma_{\text{credits}=5}(\text{Courses})) \otimes \text{Enrolled}) \otimes \text{Students})$$

A query optimizer can find this, given the first solution!

Find students enrolled at a 4 or 5 credits course

- Can identify all 4 or 5 credits courses, then find students who're enrolled in one of these courses:

$$\rho (TempCourses, (\sigma_{credits=4 \vee credits=5}(Courses)))$$
$$\pi_{name} (TempCourses \otimes Enrolled \otimes Students)$$

- Can also define TempCourses using union!
- What happens if \vee is replaced by \wedge in this query?

Find students enrolled at a 5 and 4 credits course

■ Previous approach won't work! Must identify students who're enrolled at 4 credits courses, students who're enrolled at 5 credits courses, then find the intersection (**note that *sid* is a key for Students**):

$$\rho (Temp4, \pi_{sid}(\sigma_{credits=4} (Courses) \otimes Enrolled))$$
$$\rho (Temp5, \pi_{sid}(\sigma_{credits=5} (Courses) \otimes Enrolled))$$
$$\pi_{name} ((Temp4 \cap Temp5) \otimes Students)$$

Find names of students enrolled at all courses

- Uses division; schemas of the input relations must be carefully chosen:

$$\rho (TempSIDs, \pi_{sid, cid}(Enrolled) / \pi_{cid}(Courses))$$

$$\pi_{name}(TempSIDs \otimes Students)$$

Extended Relational Algebra Operations

- Generalized Projection
- Aggregate Functions
- Outer Join
- Database modification

Generalized projection

- Extends the projection operation by allowing arithmetic functions to be used in the projection list.

$$\pi_{F_1, F_2, \dots, F_n} (R)$$

- R is any relational-algebra expression
- Each of F_1, F_2, \dots, F_n are arithmetic expressions involving constants and attributes in the schema of R .

Aggregate Functions and Operations

- **Aggregation function** takes a collection of values and returns a single value as a result.

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- **Aggregate operation** in relational algebra

$$\mathcal{G}_{G_1, G_2, \dots, G_n} \quad F_1(A_1), F_2(A_2), \dots, F_n(A_n) \quad (R)$$

- R is any relational-algebra expression
 - G_1, G_2, \dots, G_n is a list of attributes on which to group (can be empty)
 - Each F_i is an aggregate function
 - Each A_i is an attribute name

Aggregate Operation – Example

Relation
 R :

A	B	C
α	α	7
α	β	7
β	β	3
β	β	10

$g_{\text{sum}(C)}(R)$

sum(C)
27

- Result of aggregation does not have a name
 - Can use rename operation to give it a name
 - For convenience, we permit renaming as part of aggregate operation

Outer Join

■ An extension of the join operation that avoids loss of information.

- Left Outer Join 
- Right Outer Join 
- Full Outer Join 

■ Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.

■ Uses *null* values:

- *null* signifies that the value is unknown or does not exist
- All comparisons involving *null* are (roughly speaking) **false** by definition.

Modification of the Database

■ The content of the database may be modified using the following operations:

■ Deletion $R \leftarrow R - E$

■ Insertion $R \leftarrow R \cup E$

■ Updating $R \leftarrow \pi_{F_1, F_2, \dots, F_n}(R)$

■ All these operations are expressed using the assignment operator.