

Indexes

Best Practices

S5

Indexes

- An index is an on-disk structure associated with a table or view that speeds retrieval of rows from the table or view.
- Great indexing → application fast & nimble
- Poor indexing → slows entire SQL Server

```
CREATE [ UNIQUE ] [ CLUSTERED | NONCLUSTERED ]  
      INDEX index_name  
  
ON <object> ( column [ ASC | DESC ] [ ,...n ] )  
  
[ INCLUDE ( column_name [ ,...n ] ) ]  
  
[ WHERE <filter_predicate> ]  
  
[ WITH ( <relational_index_option> [ ,...n ] ) ]
```

Index Characteristics

- Clustered versus nonclustered
- Unique versus nonunique
- Single column versus multicolumn
- Ascending or descending order on the columns in the index
- Full-table versus filtered for nonclustered indexes

Clustered vs NonClustered Index

- Clustered index: sorts and stores data rows in a table, based on search key values

```
CREATE CLUSTERED INDEX Index_Name  
ON Schema.TableName(Column);
```

- NonClustered index: key values and a pointer to data in the heap/clustered index

```
CREATE INDEX Index_Name  
ON Schema.TableName(Column);
```

Clustered vs NonClustered Index

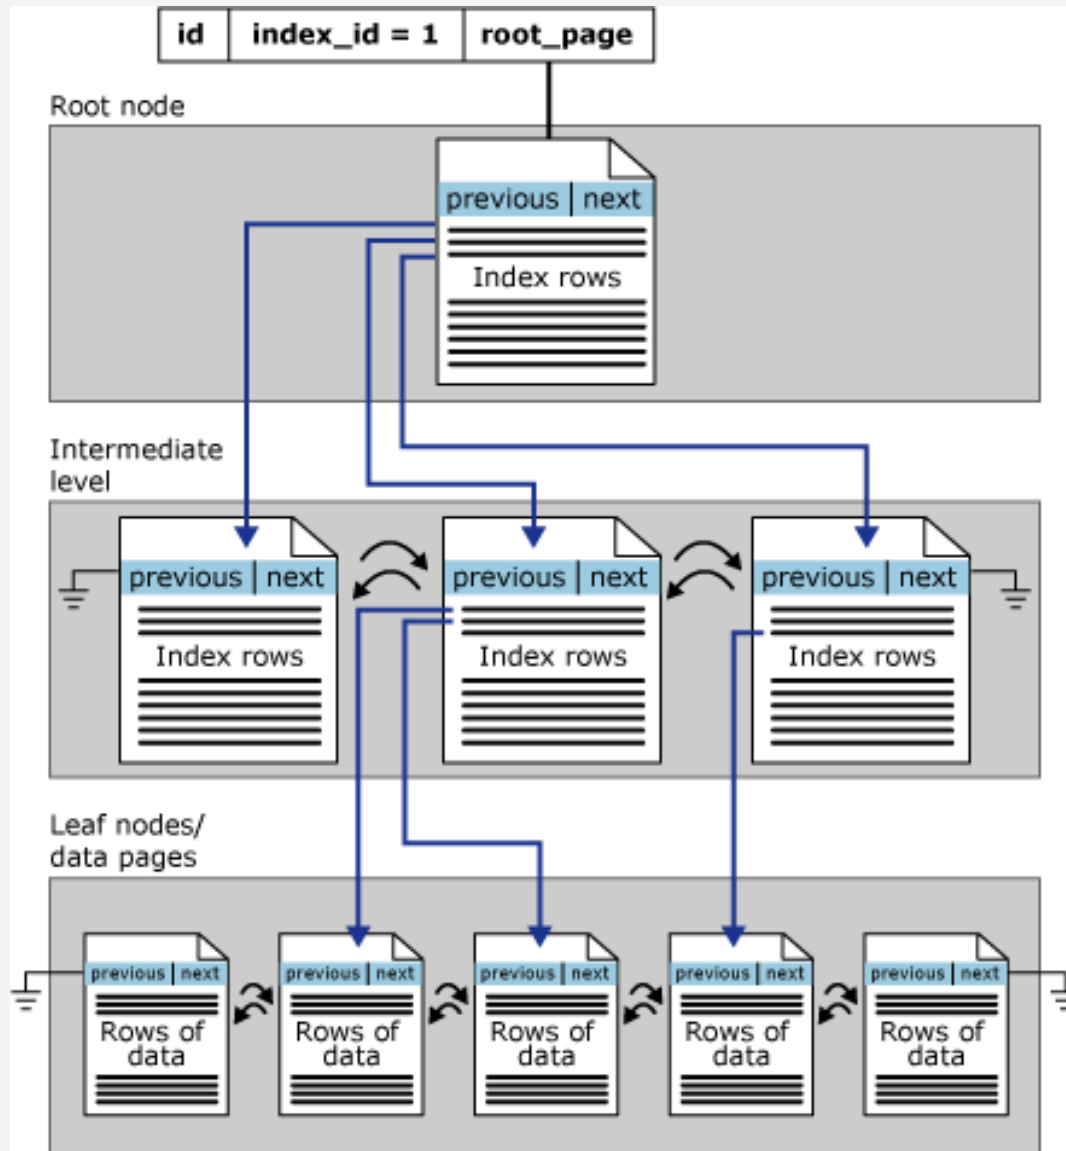
- The data pages of a clustered index will always include *all columns* in the table
- There is only one clustered index per table.
- SQL Server supports up to 999 nonclustered indexes per table.
- An index key – clustered or nonclustered – can be a maximum of 16 columns and 900 bytes.

Clustered Index

- Can be used for frequently used queries
 - Provide a high degree of uniqueness
 - Can be used in range queries
-
- Clustered indexes are not a good choice for the following attributes:
 - Columns that undergo frequent changes
 - Wide keys

Clustered Index

- indexes are organized as B-trees.



Clustered vs NonClustered Index

- When you create a primary key on a table
 - + if a clustered index is not defined
 - + a nonclustered index is not specified
 - a unique clustered index is created
- If all columns returned in a query are in the index: covering index

Key vs NonKey Index Columns

- Key columns: the columns specified to create an index.
- Nonkey columns: columns added to the INCLUDE clause of a nonclustered index.

```
CREATE INDEX Index_Name  
ON Schema.TableName(Column)  
INCLUDE (ColumnA, ColumnB);
```

Key vs NonKey Index Columns

- Benefits to using non-key columns
 - Columns can be accessed with an index scan.
 - Data types not allowed in key columns are allowed in nonkey columns (including text, ntext, and image).
 - Included columns do not count against the 900 byte index key limit enforced by SQL Server.

Index Design Tasks

- Understand the characteristics of the database
 - OLTP - On-line Transaction Processing
 - OLAP - On-line Analytical Processing
- Understand the characteristics of the most frequently used queries
- Understand the characteristics of the columns used in the queries
- Determine the optimal storage location for the index.

General Index Design Guidelines

- Database considerations
 - Too many indexes on a table affect the performance of INSERT, UPDATE, DELETE, MERGE statements
 - Indexing small tables may not be optimal
 - Indexes on views are useful when views contain aggregations and/or table joins

General Index Design Guidelines

■ Query considerations

- Create nonclustered indexes for columns frequently used in WHEREs and JOINs
- Covering indexes can improve query performance
- Write queries that insert or modify as many rows as possible in a single statement

General Index Design Guidelines

- Column Considerations
 - Keep the length of index key short for clustered indexes
 - Clustered indexes are better on unique/nonnull cols
 - Columns of **ntext**, **text**, **image**, **varchar(max)**, **nvarchar(max)**, **varbinary(max)** cannot be specified as index key columns
 - Examine column uniqueness
 - Examine data distribution in column (avoid indexes on columns with few unique values) – use filtered indexes
 - Consider the order of the columns for multiple index. Columns used in an equal to (=), greater than (>), less than (<), or BETWEEN search condition should be placed first. Additional columns should be ordered from the most distinct to the least distinct.
 - Consider indexing computed columns.

Unique Indexes

- A unique index guarantees that the index key contains no duplicate values
- Specifying a unique index makes sense only when key columns are unique
- Uniqueness – helpful information for query optimizer

Filtered Indexes

Filtered Index: an optimized nonclustered index, especially suited to cover queries that select from a well-defined subset of data

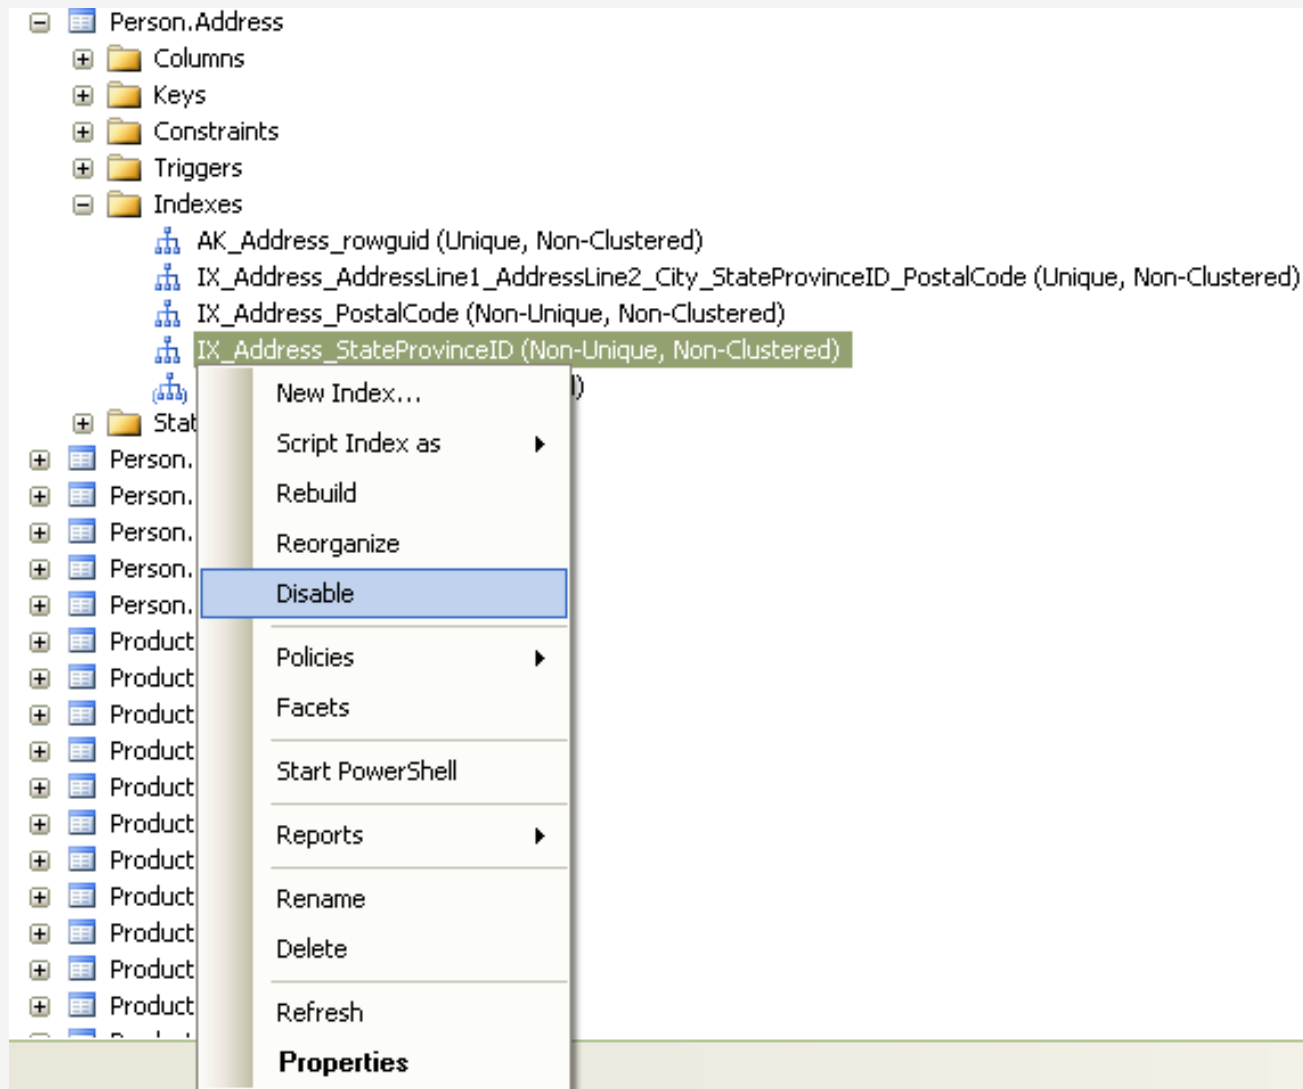
```
CREATE NONCLUSTERED INDEX FI_EndDate ON  
Products (ProductID, EndDate)  
WHERE EndDate IS NOT NULL ;  
GO
```

- Improved query performance
- Reduced index maintenance costs
- Reduced index storage costs


Disable Indexes

```
ALTER INDEX IX_Address_StateProvinceID  
ON Person.Address DISABLE
```

Dezactivarea Indecșilor



Dezactivarea Indecșilor



Enable Indexes

```
ALTER INDEX IX_Address_StateProvinceID  
ON Person.Address REBUILD
```

Indexes for Deletes

At DELETE:

- SQL Server will check for dependent rows by examining all foreign keys
- It will then check any related tables for data.
 - If there is an index, SQL Server will use that index to check for related data
 - If there isn't an index, though, SQL Server will have to **scan** the table for data.
- Deletes could be very slow if there is no index defined for foreign keys

Indexed Views

SET options

Required value

Default server value

ANSI_NULLS

ON

ON

ANSI_PADDING

ON

ON

ANSI_WARNINGS

ON

ON

ARITHABORT

ON

ON

CONCAT_NULL_YIELDS_NULL

ON

ON

NUMERIC_ROUNDABORT

OFF

OFF

QUOTED_IDENTIFIER

ON

ON

Indexed Views Restrictions

- SELECT statement cannot reference other views
- All functions must be deterministic
- AVG, MIN, MAX, STDEV, STDEVP, VAR and VARP are not allowed
- The index must be both clustered and unique
- SELECT statement must not contain subqueries, outer joins, EXCEPT, INTERSECT, TOP, UNION, ORDER BY, DISTINCT etc

Columnstore Indexes

- Groups and stores data for each column and then joins all the columns to complete the whole index
- Suited for warehouses (read only tables)
- Up to 10x query performance (vs. traditional row-oriented storage)
- Up to 10x data compression over the uncompressed data size
- The same table can have both row store index and column store index, The *Query Optimizer* will decide when to use the column store index and when to use other types of indexes

Row store for B-Tree or Heap

Row 1	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 2	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 3	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 4	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 5	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10

Page 1

Row 6	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 7	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 8	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
.....	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row n	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10

Page 2

Column Store Index

Row 1	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 2	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 3	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 4	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 5	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 6	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 7	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 8	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
.....	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row n	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
	Page 1	Page 2	Page 3	Page 4	Page 5	Page 6	Page 7	Page 8	Page 9	Page 10

Hard and fast rules for indexing

- Each table should have a clustered index that is (ideally) small, selective, ever increasing, and static. (a table without a clustered index is called a *heap*.)
- Implement nonclustered indexes on foreign key relationships
- Implement nonclustered indexes on columns that are frequently used in WHERE clauses.
- Do not implement single-column indexes on every column in a table. This will cause high overhead.
- In multi-column indexes, list the most selective (nearest to unique) first in the column list.
- For most often-used queries create covering nonclustered index.

Fragmentation

- *Internal Fragmentation*: records are stored non-contiguously inside the page. Internal fragmentation occurs if there is unused space between records in a page. The fullness of each page can vary over time. This unused space causes poor cache utilization and more I/O, which ultimately leads to poor query performance.

Fragmentation

- *External Fragmentation:* When on disk, the physical storage of pages and extents is not contiguous. When the extents of a table are not physically stored contiguously on disk, switching from one extent to another causes higher disk rotations.

Fragmentation

- *Logical Fragmentation*: Every index page is linked with previous and next page in the logical order of column data. Because of Page Split, the pages turn into *out-of-order* pages.
- An *out-of-order* page is a page for which the next physical page allocated to the index is not the page pointed to by the next-page pointer in the current leaf page.


Page read requests: 2

Extent switches: 0

Disk space used by table: 16 KB

avg_fragmentation_in_percent: 0

avg_page_space_used_in_percent: 100



Extent 1

Page 1
Data Row 1
Data Row 2

Page 6
Data Row 11
Data Row 12

Extent 2

Page 2
Data Row 3
Data Row 4

Page 4
Data Row 7
Data Row 8

Extent 3

Page 3
Data Row 5
Data Row 6

Page 5
Data Row 9
Data Row 10

Page read requests: 6

Extent switches: 5

Disk space used by table:
48 KB

avg_fragmentation_in_
percent > 80

avg_page_space_used_in_
percent: 33

Fragmentation

- *sys.dm_db_index_physical_stats*
 - **avg_fragmentation_in_percent:** This is a percentage value that represents external fragmentation.
 - **avg_page_space_used_in_percent:** This is an average percentage use of pages that represents to internal fragmentation.
- **Reducing Fragmentation in a Heap:**
 - To reduce the fragmentation of a heap, create a clustered index on the table.
 - Creating the clustered index: rearrange the records in an order, and then place the pages contiguously on disk.

Fragmentation

Reducing Fragmentation in a Index:

- If `avg_fragmentation_in_percent` $> 5\%$ and $< 30\%$, then use `ALTER INDEX REORGANIZE`:
 - reorder the leaf level pages of the index in a logical order.
- If `avg_fragmentation_in_percent` $> 30\%$, then use `ALTER INDEX REBUILD`:
 - replacement for `DBCC DBREINDEX` to rebuild the index online or offline. In such case, we can also use the drop and re-create index method.
- Drop and re-create the clustered index:
 - Re-creating a clustered index redistributes the data and results in full data pages. The level of fullness can be configured by using the `FILLFACTOR` option in `CREATE INDEX`.